

Unix Security

Processes

Adrian Coleșa

Universitatea Tehnică din Cluj-Napoca
Computer Science Department

November 2, 2015

The purpose of this lecture

- 1 presents basic concepts about Unix processes and process operations
- 2 presents specific code vulnerabilities introduced when working with processes

Outline

- 1 Processes
 - Operations on Processes
 - Program Invocation
 - Process Attributes
 - File Descriptors
 - Environment Arrays
 - Shell Variables
- 2 Interprocess Communication
 - Pipes
 - System V IPC
 - UNIX Domain Sockets

Outline

1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

Outline

1 Processes

- Operations on Processes

- Program Invocation

- Process Attributes

- File Descriptors

- Environment Arrays

- Shell Variables

2 Interprocess Communication

- Pipes

- System V IPC

- UNIX Domain Sockets

Process Creation

- using the *fork()* system call
- the child is an identical copy of its parent
 - they share: memory, attributes, open files
 - the new process is given a new PID (process ID)
 - both processes continue their execution with the first instruction after *fork*
 - *fork* returns child PID in parent and 0 in child
- the parent-child relationship is tracked by the OS
- the child can get its parent's ID with *getppid()*
- processes whose parent terminates before them are given as their new parent the *init* process

General Parent-Child Template

```
pid_t pid;

switch (pid = fork()) {
    case -1:
        perror("fork");
        break;
    case 0:
        printf("Child: pid=getpid(), ppid=getppid()\n");
        // ... some other child job
        exit(0);
    default:
        printf("Parent: pid=getpid(), ppid=getppid()\n");
        // ... some other parent job
        exit(0);
}
```

fork() Variants

- there are some other variants of the classic fork
- *vfork()* use to avoid the performance paid for coping memory for a new process that immediately loads a new code
 - memory was shared
 - the parent is blocked until child loads other code or terminates
 - shared memory is supposed not to be changed by the child
 - as copy-on-write become common, it gets deprecated
- *rfork()* from plan9 OS is used to let the user specify the shared resources at a more granular level
- *clone()* is a similar correspondent in Linux

fork() Variants (cont.)

- they are used mainly for thread (lightweight processes) creation

Process Termination

- several ways (and reasons)
- voluntarily using *exit()*
- involuntarily, being terminated by the system by sending them signals
 - reasons: exceptions, processes sending signals, abort
 - default handling of signals is to terminate the process
 - though, some signals could be explicitly handled by the process

fork() and Open Files

- OS open file management tables
 - process file descriptor table (FDT)
 - system open file table (OFT)
 - system i-node table (IT)
- child inherits file descriptors from its parent ⇒
 - both parent's and child's FDs reference the same entries in the system OFT
 - open files are shared between parent and child ⇒
 - **possible race conditions**
- file descriptors can evolve independently in parent and child after fork

Outline

1 Processes

- Operations on Processes
- **Program Invocation**
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

Direct Invocation

- using one function of the *exec* family
 - *execl()*, *execlp()*, *execle()*, *execv()*, *execv1()*
- the most generic one: *execve()*

```
int execve(const char *path, char *const argv[], char *const envp[]);
```

- *argv*: command line
- *envp*: environment variables

Dangerous *execve()* Variants

- all function share the same security issues with *execve()*
- *exec/p()* and *execvp* have additional concerns
- they are based on the value of PATH environment variable
- if the attacker could control the PATH, he could control which program to be loaded

Dangerous `execve()` Variants (cont.)

```
int print_directory_listing(char *path)
{
    char *argv[] = {"ls", "-l", path, NULL};
    int rc;

    rc = fork();

    if (rc < 0)
        return -1;

    if (rc == 0)
        execvp("ls", argv);
    return 0;
}
```

- setting for instance `PATH` to `"/tmp"` will make running a program `"/tmp/ls"`

The Argument Array

- programs usually use *switch* instruction to process their user received arguments
- some programs fail to sanitize their arguments correctly
- example (from *vacation* program): vulnerable code not sanitizing user supplied argument, which could influence the called *sendmail* program

```
void sendmessage(char *myname)
{
    ...
    if (vfork() == 0) {
        execlp(_PATH_SENDMAIL, "sendmail", "-f", myname, from, NULL);
    }
    ...
}
```


The Argument Array (cont.)

- attack
 - send email from address “-C/some/file/here”
 - control *sendmail* to load an alternative configuration file
 - ⇒ execute arbitrary commands on behalf of the vacationing user
- code audit: when a program use `getopt` function be aware of
 - if the program considers option arguments in the same string with the option (like “-C/some/file”)
 - after “--” the options are considered normal arguments and not handled by *getopt*

Indirect Invocation. Overview

- using functions that run a sub-shell
- specify a command line interpreted by a sub-shell
- popular functions
 - C: *system()*, *popen()*
 - Perl: *system*, *open()*
 - Java: *Runtime.getRuntime().exec()*
 - Python, PHP etc.

Indirect Invocation. Security Problems

- meta-characters
 - command separators, file redirection, evaluation operators
- globbing chars for FS access
 - wildcards used to locate files based on a pattern: “.,?*[]{}”
 - inherent in shell interpreters
- environment issues
 - shell tends to change their functionality based on certain environment variables
- SUID Shell Scripts
 - generally a bad idea → can easily be tricked by meta-characters and globbing

Outline

1 Processes

- Operations on Processes
- Program Invocation
- **Process Attributes**
- File Descriptors
- Environment Arrays
- Shell Variables

2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

Process Attributes Retention

- when loading a new program (with *exec*), memory is remapped
- there are attributes inherited by the new program
- can be sources of potential vulnerabilities when
 - new application is more privileged
 - current application drops its privileges by loading the new one
- inherited attributed
 - file descriptors
 - signal mask, even if signal handlers are lost
 - effective UID / GID, except the case the new program is SUID / SGID

Process Attributes Retention (cont.)

- real UID / GID
- PID, PPID and process group ID
- supplemental groups
- working and root directory
- controlling terminal
- resource limits
- umask

Resource Limits (*rlimits*)

- enforce restrictions on the system resources that a process may use
- functions to manipulate limits: *getrlimit()* and *setrlimit()*
- each resource has two limits associated to: *soft* and *hard*
- examples of resource limits for a process
 - RLIMIT_CORE: maximum size for a core file
 - RLIMIT_CPU: maximum CPU time (sec)
 - RLIMIT_DATA: maximum size (bytes) for the data segment
 - RLIMIT_FSIZE: maximum size of a written file
 - RLIMIT_MEMLOCK: maximum no of bytes locked in memory
 - RLIMIT_NOFILE: maximum number of open files

Resource Limits (*rlimits*) (cont.)

- RLIMIT_NPROC: maximum no of processes a user can run
- RLIMIT_STACK - maximum size (bytes) for process' stack
- rlimits are useful to restrict a process

Resource Limits Vulnerabilities

- security risks
 - rlimits settings survive the *exec* calls
- attack method
 - force a called privileged process to fail in a predetermined location
- caused by
 - rlimits overrun errors not handled appropriately
 - e.g. signaled not handled
 - e.g. unfounded trusted in environment (“improbable conditions”)

Resource Limits Vulnerabilities. Example 1

```
if (!(found = !uselib(buff))) {  
    if (errno != ENOENT) {  
        fdprintf("2, %s: cannot load library '%s'\n", argv0, buff);  
    }  
}
```

- vulnerable due to
 - *buffer-overflow* in *fdprintf*
 - triggered by *overrunning rlimits*
- attack vector
 - exhaust all applications file descriptors
 - provide a long special crafted application name (argv[0])

Resource Limits Vulnerabilities. Example 2

- privileged code vulnerable to not handled `RLIMIT_FSIZE`

```
struct entry {
    char name[32];
    char password[256];
    struct entry *next;
};

int write_entries(FILE *fp, struct entry *list)
{
    struct entry *ent;

    for (ent = list; ent; ent=ent->next)
        fprintf(fp, "%s:%s\n", ent->name, ent->password);
    return 1;
}
```

Resource Limits Vulnerabilities. Example 2 (cont.)

- attack vector
 - set a low `RLIMIT_FSIZE`
 - mask signal `SIGXFSZ` (to be ignored) before calling the privileged program
 - could cause partial writing, e.g. truncating a password

Resource Limits Vulnerabilities. Code Audit

- check for write operations, whose result is not checked
 - both success/fail and no of written bytes
- never assume that a condition is unreachable because it seems unlikely to occur
 - rlimits could trigger such conditions by restricting resources of a privileged program

Outline

1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- **File Descriptors**
- Environment Arrays
- Shell Variables

2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

File Sharing Aspects

- due to *fork* or *dup* multiple fd across one or more processes refer the same open file object
 - they share all open file properties, like open mode, current position
- when multiple processes open the same file,
 - they share the same physical file and properties
 - each have a different logical view, i.e. its own open file (access mode and current position)
- an open file object and corresponding access is kept unchanged by a process even if
 - owner changed
 - permissions changed
 - file (path) removed

Close-on-Exec

- file descriptors retained over *execve()*
- UNLESS explicitly marked for closure
- setting file descriptors for *close-on-exec*
 - is a useful precaution for sensitive files to not be inherited by a subprogram
 - can be done at *open()* or with *fcntl()*
- code audit
 - for applications that creates new processes
 - check to see if there are opened files not marked for close and evaluate their implication

File Descriptor Leaks. Overview

- context
 - security checks done only at the opening of a file
 - access kept even if permissions are restricted or the application privileges are lost
- risks
 - new loaded (user controlled) code could use unintentionally inherited file descriptors
- recommendation
 - programs working with file descriptors to security-sensitive resources, should close their descriptors

File Descriptor Leaks. Examples

- vulnerable code that does not close a file descriptor to (device driver to) kernel memory

```
int kfd;
pid_t p;
char *initprog;

kfd = safe_open("/dev/kmem", O_RDWR);
init_video_mem(kfd);

if (initprog = getenv("CONTROLLER_INIT_PROGRAM")) {
    if (p=safe_fork()) { //parent
        wait_for_kid(p);
        g_controller_status = CONTROLER_READY;
    } else { //child
        drop_privs();
        execl(initprog, "conf", NULL);
        exit(0);
    }
}
```

File Descriptor Leaks. Examples (cont.)

- similar real vulnerabilities: *libkvm* (FreeBSD), *chpass* (OpenBSD)

File Descriptor Leaks. Code Audit

- programs that drop privileges to run unsafe code
 - should be evaluate from the perspective of file descriptor management
- not limited just to files
 - any resource that can be represented with a file descriptor: pipes, sockets, etc.

File Descriptor Omission. Overview

- FD allocation → the lowest available
- special (system) FDs
 - 0: STDIN
 - 1: STDOUT
 - 2: STDERR
- certain library functions consider the default associations
 - *scanf*, *gets* → `read(0, ...)` //STDIN
 - *printf*, *puts* → `write(1, ...)` //STDOUT
 - *perror* → `write(2, ...)` //STDERR
- privileged programs could be tricked to
 - write sensitive data into attacker's files
 - get inputs from attacker's files

File Descriptor Omission. Overview (cont.)

- attack vector
 - starts a SUID program with the standard file descriptors closed
 - any new file the program will open will be allocated one of the standard descriptors
 - \Rightarrow program could leak important output to attacker
- fixes of such vulnerabilities
 - checking if 0, 1, and 2 are available and allocate them for `"/dev/null"`

File Descriptor Omission. Example

- vulnerable due to the possible allocation $fd = 2$

```
if ((fd = open("/etc/shadow", O_RDWR)) < 0)
    exit(1);

user = argv[1];

if ((id = find_user(fd, user)) < 0) {
    fprintf(stderr, "Error: invalid user %s\n", user);
    exit(1);
}
```

- attack

- `close(2)` \Rightarrow the error message will go into `"/etc/shadow"`
- username contains `'\n'` to introduce new user accounts (with root permissions) into `"/etc/shadow"`

Outline

1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- **Environment Arrays**
- Shell Variables

2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

Process Environment

- a collection of pairs of type “NAME=VALUE”
- a process’ environment is maintained by the standard library
 - represented in memory as an array (*environ*) of pointers to C-like strings
 - last element is NULL
- `execve` passes the new program its environment
 - the kernel copies the environment variables into the memory of the new program at adjacent locations
 - would likely be next to program argument strings at the top of the program’s stack

Process Environment (cont.)

- as a process run, it can add, modify or delete its environment variables
- new environment strings are allocated in the heap with *malloc*
- functions to manipulate the environment: *getenv*, *setenv()*, *putenv()*, *unsetenv*, *clearenv*
- the standard C library expects the strings to be in the particular format (two strings separated by '=')

Confusing *putenv()* and *setenv()*

- *putenv()* not make a copy of the string passed as argument
 - inserts the pointer directly into the environment array
 - \Rightarrow user can later modify data that is pointed by that pointer
 - \Rightarrow the pointer could be discarded
- vulnerable code: after function returns the pointer points to undefined stack data!

```
int set_editor(char *editor)
{
    char edstring[1024];

    snprintf(edstring, sizeof(edstring), "EDITOR=%s", editor);

    return putenv(edstring);
}
```

Extraneous Delimiters

- variable name or value contain an extra '='
- old library functions (e.g. *setenv*, *unsetenv*) dealt differently with such cases
- current implementations normally do not accept having '=' inside the variable name
- care must be taken when an application has its own implementation for environment variable management
 - take a look at and compare how variables are found and set
 - take a look at code making assumptions the variable name contains no special delimiters

Extraneous Delimiters (cont.)

- vulnerable code: allows addition of a variable with an arbitrary value

```
int set_var(char *name)
{
    char *newenv;

    int len = strlen("APP_") + strlen("=new") + strlen("name") + 1;
    newenv = (char*) malloc(len);
    snprintf(newenv, len, "APP_%s=new", name);
    return putenv(newenv);
}
```

Duplicate Environment Variables

- have more variables with the same name defined in the environment
- current library functions are safe from this point of view
- code review: look at custom implementations
- vulnerable code: missing two consecutive entries with the same name

Duplicate Environment Variables (cont.)

```
static void _dl_undestenv(const char *var, char **env)
{
    char *ep;

    while ((ep = *env)) {
        const char *vp = var;

        while (*vp && *vp == *ep) {
            vp++;
            ep++;
        }
        if (*vp == '\0' && *ep++ == '=') {
            char **P;

            for (P=env; ++P)
                if (!(*P == *(P + 1)))
                    break;
        }
        env++;
    }
}
```

Shellshock Vulnerability

- September 2014: CVE-2014-6271, CVE-2014-6277, CVE-2014-6278, CVE-2014-7169
 - affected versions Bash up to 4.3
 - vulnerability overview
 - environment variables whose value started with '()' interpreted as function definitions
 - not correctly treated: allowed extra commands to be specified after the normal function definition
 - exploitation example
- ```
env x='()' { :;; echo vulnerable' bash -c "echo this is a test"
```
- specific exploitation methods



## Shellshock Vulnerability (cont.)

- CGI-based web server
- OpenSSH server
- DHCP clients

# Outline

## 1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- **Shell Variables**

## 2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

# PATH Environment Variable

- a list of directories separated by ':'
- an executable (command) name is searched in those directories
- current directory is searched only if specified explicitly
- vulnerable code: when part of a privileged application  
"/opt/ttt/start\_process"

```
snprintf(buf, sizeof(buf),
 "/opt/ttt/logcat %s | gzip | /opt/ttt/parse > /opt/ttt/results", logfile);
system(buf);
```

## PATH Environment Variable (cont.)

- attack vector

```
$ cd /tmp
$ echo '#!/bin/sh' > gzcat
$ echo 'cp /bin/sh /tmp/sh' >> gzcat
$ echo 'chown root /tmp/sh' >> gzcat
$ echo 'chmod 4755 /tmp/sh' >> gzcat
$ chmod 755 ./gzcat
$ export PATH=.:$PATH
$ /opt/ttt/start_process
$ /tmp/sh
#
```

# HOME

- indicate where the user's home directory is placed in the file system
- used in cases like "`~/file`"
- an attacker (user) can change the variable, so it's good for a privileged application to check the path also in password database

# IFS

- IFS = internal field separator
- tells the shell which characters represent white spaces (normally spaces, tabs, and new lines)
- if an attacker changes the IFS, it could run privileged code
- vulnerable code: not check or set IFS

```
system("/bin/ls");
```

## IFS (cont.)

- attack vector

```
$ cd /tmp
$ echo 'sh -i' > bin
$ chmod 755 ./bin
$ export PATH=.:$PATH
$ export IFS='/'
$ vuln_program # "/bin/ls" interpreted like "bin ls"
$./sh
#
```

- normally not working on modern shells that filters dangerous environment variables like IFS

## Other Dangerous Environment Variables

- ENV (or BASH\_ENV)
  - used by a non-interactive shell to run the associated filename as a startup script
  - it is usually expanded
  - attack example: `ENV=``/tmp/evil```
  - any subshells that are opened actually run the `“/tmp/evil”`
- SHELL
  - indicate the user preferred shell
- EDITOR
  - indicate the user preferred editor



# Runtime Linking and Loading Variables

## • LD\_PRELOAD

- provides a list of libraries that the runtime link editor loads before it loads everything else
- gives a chance to the user to insert his own code into a process or his choosing
- in general UNIX OSes do not honor `LD_PRELOAD` when running SUID and SGID programs

## • LD\_LIBRARY\_PATH

- provides a list of directories containing shared libraries
- the runtime link editor searches through this list first when looking for shared libraries
- it is ignored for SUID/SGID binaries

# Outline

## 1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

## 2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

# Outline

## 1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

## 2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

# Anonymous Pipes

- a uni-directional pair of file descriptors
- one file descriptor for read, one for write
- created and automatically opened using *pipe()* system call
- the underlying mechanism used to run commands like

“`cmd1 | cmd2`”

- the shell creates an anonymous pipe
- the shell creates two processes
- the first process' STDOUT is redirected to the write file descriptor of the pipe
- the second process' STDIN is redirected to the read file descriptor of the pipe

## Anonymous Pipes (cont.)

- first process runs first command, the second runs the second command
- *popen* system call
- writing to a pipe with no read file descriptor causes the writing program to receive a SIGPIPE signal

## Named Pipes (FIFO files)

- have a name and can be opened like any other normal file
- created using *mkfifo* or *mknod*, opened with *open*
  - vulnerable at race conditions attacks, because they only create a FIFO file, but not open it
  - between creation and opening the FIFO could be replaced by an attacker
  - example: vulnerable code

```
int open_pipe(char *pipename)
{
 int rc;

 if ((rc = mkfifo(pipename, S_IRWXU)) < 0)
 return -1;
 return open(pipename, O_WRONLY);
}
```

## Named Pipes (FIFO files) (cont.)

- *open* could be blocking, if mode is just for read or write only, until a counter peer process occurs
  - not a security problem in-itself, it could be used as a slowing-down (even blocking) a process in a TOCTOU attack
  - exploitation: an application opening a regular file is provided a named pipe
  - application could even be finer-tune controlled, if the attacker's application is the only writer at the other end of the pipe
- non-blocking behavior could be set explicitly
- code audit should check for the implications of

## Named Pipes (FIFO files) (cont.)

- pipes created with insufficient privileges giving an attacker access to the pipe and interfering with the normal IPC
- applications intending to work with a regular file, but being provided a FIFO, because failing to determine the file type
- race conditions introduced by *mkfifo* (*mknod*) and *open*



# Outline

## 1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

## 2 Interprocess Communication

- Pipes
- **System V IPC**
- UNIX Domain Sockets

# Message Queues and Shared Memory

- message queues
  - a simple stateless messaging system
  - a sort of specialized message-based pipes
  - unlike file system access, message queue permissions are checked for each operation
  - functions: *msgget()*, *msgctl()*, *msgrcv()*, and *msgsnd()*
- shared memory
  - the mechanism that maps the same memory segment to more processes' address spaces
  - functions: *shmget()*, *shmctl()*, *shmat()*, and *shmdt()*
- have their own namespace in kernel memory, not tied to the FS

## Message Queues and Shared Memory (cont.)

- implement their own simple permissions model
- code audit: check for improper permissions of System V IPC
- after a process fork both parent and child have a copy of the mapped shared memory
- after `exec()` the shared memory is detached
- use of shared resources could introduce risks of race conditions
  - ex.: if multiple processes share the same memory, one process could change (write) some data after a process has just read (check) it

# Outline

## 1 Processes

- Operations on Processes
- Program Invocation
- Process Attributes
- File Descriptors
- Environment Arrays
- Shell Variables

## 2 Interprocess Communication

- Pipes
- System V IPC
- UNIX Domain Sockets

# UNIX Domain Sockets

- similar to pipes (also anonymous and named)
- allow local processes to communicate each other
- anonymous domain sockets are created by using *socketpair()* function
  - creates a pair of unnamed endpoints that a process can use to communicate information its next children
- named sockets use the socket API functions - similar to networked applications
  - implemented using special socket device files, created automatically when a server calls *bind()*
  - location of the filename is specified in the socket address structure

## UNIX Domain Sockets (cont.)

- created with permissions “`777 & ~umask`” - exposed to attacks if “`umask = 0`”
- vulnerable code: not setting *umask* before creating the socket

```
int create_sock(char *path)
{
 struct sockaddr_un sun;
 int s;

 bzero(&sun, sizeof(sun));
 sun.sun_family = AF_UNIX;
 strncpy(sun.sun_path, path, sizeof(sun.sun_path) - 1);

 if ((s = socket(AF_UNIX, SOCK_STREAM, 0)) < 0)
 return -1;

 if (bind(s, (struct sockaddr) (&sun, sizeof(sun))) < 0)
 return -1;
 return s;
}
```

## UNIX Domain Sockets (cont.)

- code is also vulnerable to race conditions
  - if the user can specify parts on entirely the socket pathname
  - e.g. if user writable directories are used in the path

# Bibliography

- 1 “The Art of Software Security Assessments”, chapter 10,  
“Strings and Metacharacters”, pp. 559 – 624